

Decision Trees and Pegasos SVMs from Scratch: Price and Price-Tier Prediction on bina.az Listings

Surname1, Name1 Surname2, Name2 Surname3, Name3

MLE-AI-201 Machine Learning, AI Academy — Cohort II

Abstract—We implement a CART decision tree (Gini, entropy and MSE criteria) and a soft-margin support-vector machine trained with the Pegasos sub-gradient method entirely in NumPy, and apply them to 57,721 cleaned sale listings scraped from bina.az, Azerbaijan’s largest real-estate portal. One dataset is framed as two tasks: regression of $\log(\text{price})$ (Task A) and classification of a *premium* vs. *standard* price tier split at the training median (Task B). Every hyperparameter is selected on a validation split; final models are refitted on train+validation and scored once on a held-out test set, with 5-fold cross-validation for uncertainty. Our tree reaches a test RMSE of 0.229 on $\log(\text{price})$ ($R^2 = 0.866$), against 0.229 for scikit-learn’s tree with matched settings; a node-by-node comparison leaves 0 unexplained differences, with the remaining ones due to exact gain ties and scikit-learn’s float32 cast of the features. Our averaged Pegasos SVM reaches test F1 0.912 (scikit-learn LinearSVC: 0.911), and a random-Fourier-feature extension reaches 0.914. Error analysis shows that misclassified tiers concentrate around the threshold, and that the regression tree under-predicts the most expensive listings. Code and one-command reproduction: github.com/eyyubovsamur240-afk/ML-final-project (tag v1.0-final).

Index Terms—decision tree, support-vector machine, Pegasos, random Fourier features, real-estate price prediction.

I. INTRODUCTION

Asking prices on listing portals are noisy, heavy-tailed and driven by location, size and condition. We use them to study two classic learners from the inside: we write both models ourselves, feed them a real, uncleaned scrape, and benchmark them honestly against scikit-learn. Our contributions are:

- a leakage-free pipeline for the bina.az dump: tolerant parsing of Azerbaijani, unit-laden fields, currency conversion, documented cleaning rules, and preprocessing fitted on training rows only;
- a vectorised CART tree whose split search evaluates every threshold of every feature with cumulative sums, verified to make the same split as scikit-learn at every node up to exact ties;
- a Pegasos SVM with mini-batches, projection, iterate averaging and a random-Fourier-feature (RFF) kernel extension, whose objective we compare with LIBLINEAR’s exact optimum;
- a protocol that never touches the test split during selection, CV uncertainty for all models, error analysis, and bonus ensembles, ridge regression and clustering built on our own code.

II. BACKGROUND

Decision trees. CART [2] grows a binary tree greedily: each node picks the feature/threshold pair that maximises the impurity decrease $\Delta I = I(\text{parent}) - \sum_c \frac{N_c}{N} I(c)$, where I is the Gini index $1 - \sum_k p_k^2$, the entropy $-\sum_k p_k \log_2 p_k$ (information gain [3]), or, for regression, the node variance $\frac{1}{N} \sum_i (y_i - \bar{y})^2$. Growth stops on depth, node-size or minimum-gain rules; leaves predict the majority class (with class frequencies as probabilities) or the mean.

SVMs and Pegasos. The soft-margin SVM [4] minimises the regularised hinge loss. Pegasos [1] solves this

primal by stochastic sub-gradient descent with step $\eta_t = 1/(\lambda t)$ and reaches ϵ accuracy in $\tilde{O}(1/(\lambda \epsilon))$ steps, independent of n . Random Fourier features [5] map $\mathbf{x}$ to $\phi(\mathbf{x}) = \sqrt{2/D} \cos(W\mathbf{x} + \mathbf{c})$ with $W \sim \mathcal{N}(0, 2\gamma I)$, so that $\phi(\mathbf{x})^\top \phi(\mathbf{z}) \approx e^{-\gamma \|\mathbf{x} - \mathbf{z}\|^2}$ and a linear solver on ϕ approximates an RBF-kernel SVM.

III. DATA AND PREPROCESSING

A. Dataset and cleaning

We use the Kaggle *Bina.az sale project* dump [13] (file `house_sale.csv`, SHA-256 prefix `b6a4c67e0d3c`): 100,775 rows scraped between 2024-10-05 and 2024-11-19. Column names are Azerbaijani (*Sahə, Otaq sayı, Mərtəbə, ...*), most numbers arrive as strings with units (“145 m²”, “7 / 9”, “2 640 AZN/m²”) and every price carries a currency field (AZN 64,454). All cleaning lives in `src/data_prep.py`:

1. names are transliterated and mapped to English (*area, rooms, floor, land area, repair, mortgage, bill of sale, category*); where the dump holds two versions of a field (repair as “Təmirli”/blank and as *var/yoxdur*) the complete one is used;
2. tolerant parsers handle thousands separators, decimal commas, “floor / total floors”, land in *sot*/m²/ha and yes/no flags; prices are converted to AZN at fixed rates;
3. the scraper visited many listings several times, sometimes after a price change: the 100,775 rows hold only 64,454 distinct listing URLs, and we keep the latest scrape of each;
4. we study *residential* listings (new builds, old stock, houses/villas; 58,452 listings). Land plots (area in *sot*, no rooms or floors), commercial objects, offices and garages (torpaq 3,315, obyekt 2,526, ofis 101, qaraj 60) are different markets; the scope is one setting in `config.py`;

Table 1. Cleaning steps and rows remaining.

Step	Rows left	Removed
raw rows	100,775	
repeat scrapes of a listing removed (latest kept)	64,454	36,321
residential categories only	58,452	6,002
has a positive price	58,452	0
has a positive area	58,452	0
exact duplicates removed	58,110	342
re-posted listings removed	57,737	373
price in [5,000, 15,000,000] AZN	57,728	9
area in [10, 3000] m ²	57,724	4
price/m ² in [30, 25,000] (typo filter only)	57,721	3
final cleaned rows	57,721	0

- rows without a positive price or area are dropped, then exact duplicates and re-posts of the same flat under a new listing id (same price, area, rooms, floor and coordinates);
- fixed plausibility rules remove typos and unit mix-ups (Table 1); impossible feature values (floor above the building height, coordinates outside Azerbaijan) are set to missing and imputed later.

The rules are constants chosen from domain knowledge, not statistics estimated on the data, so applying them before the split cannot leak.

Leakage. We drop **unit_price**, **total_price** because they are arithmetic functions of the target (price per m², price restated); the cleaning code also drops any other column whose name contains “price”. We also drop identifiers and free-form addresses (id_x, rel_url, img_url, id_y, estate_id, estate_rel_url_y, owner_name, owner_title, shop_name, shop_title, address, estate_details_id_x, estate_details_id_y, estate_rel_url, listing_key) and platform promotion flags (scraped_at, city_when, day_x, hour_x, vip, featured, products_label, datetime_scrape_y, currency_y, updated, views, day_y, hour_y), which describe the seller’s advertising budget rather than the property. Price per m² is used once, only as a typo filter, and never as a feature.

B. Features, tasks and split

From the cleaned rows we build 24 numeric/binary features (area and log area, rooms, area per room, floor, total floors, floor ratio, top/ground-floor flags, land area, latitude, longitude, distance to the Baku centre, repair/mortgage/bill-of-sale flags and six keyword flags from the description, e.g. “metro” or “dəniz” (sea)), plus one-hot encodings of category, building type, city, administrative district (parsed from bina.az’s location tags) and the finer location label (metro station or settlement), 113 columns in total. One-hot rather than ordinal encoding is used because none of these has a natural order. Missing numeric values are imputed with the *training* median and flagged by an indicator, and a category level gets its own column only if it has at least 20 *training* rows; rarer or unseen levels fall into an “other” column.

Task A predicts log(price): the raw price is strongly right-skewed (skewness 6.60; 0.32 after the log, Fig. 1a–b), and an error in log space is a relative error. Task B labels a listing *premium* if its price exceeds the *training* median, 208,000 AZN. The split is 70/15/15 (40,405/8,658/8,658), stratified

Table 2. Dataset summary after cleaning.

Property	Value
Rows (raw / cleaned)	100,775 / 57,721
Features after encoding	113
Train / Val / Test	40,405 / 8,658 / 8,658
Tier threshold (train median)	208,000 AZN
Premium share train / val / test	50.0% / 50.1% / 50.0%
Median price (train)	208,000 AZN
Skewness price / log(price)	6.60 / 0.32

on deciles of log(price) with a fixed seed, so all splits share the price distribution and the tier balance (50.0%, 50.1%, 50.0% premium). Stratifying only decides which rows go where; no statistic is computed from validation or test rows.

Exploration (training split only, Fig. 1 and the appendix). The median price is 208,000 AZN (1st–99th percentile 46,000–1,250,000), the median price per m² is 2,235 AZN, and log area correlates with log price at 0.72. The median price per m² varies several-fold between districts (Fig. 7), so location carries a large part of the signal. After cleaning the most frequent missing fields are building_type 100.0%, land_area_sot 83.6%, mortgage 65.8%, floor 16.4%, total_floors 16.4%.

IV. METHOD

A. Decision tree (from scratch)

DecisionTree (src/decision_tree.py) builds Node objects recursively and exposes a scikit-learn-style fit/predict/predict_proba API. *Split search:* for a node with m rows we sort all candidate features at once (one argsort of the $m \times d$ block) and evaluate *every* threshold with cumulative sums. For classification, running class counts give the Gini or entropy of both children at every cut. For regression, running $\sum y$ and $\sum y^2$ give the children’s sums of squared errors, computed after centring y at the node mean for numerical stability. Candidate thresholds are midpoints between consecutive *distinct* sorted values. A split is valid only if both children keep at least `min_samples_leaf` rows, and the best gain wins (ties go to the lowest feature index). One node therefore costs $O(dm \log m)$, almost all of it inside NumPy.

Stopping and regularisation. A node becomes a leaf if it is pure, reaches `max_depth`, has fewer than `min_samples_split` rows, has no valid cut, or if the weighted gain $\frac{N}{N} \Delta I$ is below `min_impurity_decrease`. This is the same weighted definition scikit-learn uses, so matched settings mean the same thing in both libraries. Leaves store class frequencies (giving `predict_proba`) or the mean target. Feature importance is the total weighted gain per feature, normalised to one (mean decrease in impurity). After fitting, the tree is flattened into arrays so that prediction routes all rows through the tree together.

Depth curves from a single fit. Greedy CART decides each node from that node’s rows only, so a tree grown to depth D and *truncated* at depth d is identical to a tree grown with `max_depth=d`. A unit test checks this equivalence, and we use it to draw depth curves at the cost of one fit.




Figure 1. Exploration on the training split. (a, b) Price is heavy-tailed; $\log(\text{price})$ is close to symmetric. (c) Listing prices across Baku: the centre and the seafront are the most expensive. (d) $\log$ area vs. $\log$ price (correlation 0.72).

B. Support-vector machine (from scratch)

PegasosSVM (src/svm.py) minimises

$$J(\mathbf{w}) = \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i \mathbf{w}^\top \mathbf{x}_i), \quad y_i \in \{-1, +1\}, \quad (1)$$

where $\mathbf{x}_i$ is augmented with a constant 1 so that the bias is the last weight. With a random mini-batch A_t of size k and the set A_t^+ of its hinge-active points ($y_i \mathbf{w}^\top \mathbf{x}_i < 1$), one step is

$$\mathbf{w} \leftarrow (1 - \eta_t \lambda) \mathbf{w} + \frac{\eta_t}{k} \sum_{i \in A_t^+} y_i \mathbf{x}_i, \quad \eta_t = \frac{1}{\lambda t}, \quad (2)$$

followed by the optional projection onto the ball $\|\mathbf{w}\| \leq 1/\sqrt{\lambda}$, which contains the optimum. $k = 1$ is classic Pegasos and $k = n$ is full-batch sub-gradient descent; the $\eta_0/\sqrt{t}$ and constant schedules are also implemented. Each epoch reshuffles the rows with a seeded generator and records $J(\mathbf{w})$ on the full training set. Because the last iterate is noisy when λ is small, the tuned solver returns the t -weighted iterate average $\bar{\mathbf{w}}_t = (1 - \frac{2}{t+1})\bar{\mathbf{w}}_{t-1} + \frac{2}{t+1}\mathbf{w}_t$, which keeps the optimal $O(1/(\lambda T))$ rate [6]. Features are standardised with training statistics, since the hinge loss is scale-sensitive. We report $\|\mathbf{w}\|$, the margin width $2/\|\mathbf{w}\|$, and the support vectors, defined as the training points with $y_i f(\mathbf{x}_i) \leq 1$.

Non-linear extension: random Fourier features. We map the standardised features with $D = 1,024$ random features and run the *same* linear solver on $\phi(\mathbf{x})$ (RFFPegasosSVM). We chose RFF over kernelised Pegasos because training and prediction cost $O(nD)$ instead of growing with the number

of support vectors, which matters for $\sim 49,063$ training rows, and because it reuses the tested linear code unchanged. The cost is a Monte-Carlo kernel approximation, whose error we measure (Fig. 3, bottom-c).

V. EXPERIMENTAL SETUP

Protocol. (1) All hyperparameters are chosen on the validation split with models trained on the training split; the choice rule is fixed in code (best validation score, ties to the simpler model). (2) The chosen configurations are refitted on train+validation (49,063 rows) with imputation, encoding, scaling and the tier threshold re-learned on those rows (threshold 208,000 AZN), and scored once on the test split with 95% percentile-bootstrap intervals (1,000 resamples). (3) For uncertainty, 5-fold stratified CV on train+validation repeats the whole preprocessing inside every fold and reports $\text{mean} \pm \text{std}$.

Metrics. Task A: RMSE, MAE and R^2 on $\log(\text{price})$, plus MAE (AZN) and MAPE after back-transforming with exp, which predicts the conditional median. Task B: accuracy, precision, recall and F1 of the premium class, ROC-AUC and PR-AUC (average precision) from continuous scores (leaf frequencies for trees, margins for SVMs), and confusion matrices. All metrics are our own NumPy code and are unit-tested against scikit-learn.

Baselines with matched settings. Our tree vs. DecisionTreeRegressor/Classifier with the same criterion, max_depth, min_samples_leaf and min_impurity_decrease. Our SVM vs. SGDCClassifier (hinge, $\alpha = \lambda$, “optimal” schedule [7]) and LinearSVC

Table 3. Linear Pegasos SVM vs. λ (train→validation; the selected row is bold). $C = 1/(n\lambda)$.

λ	C	$\ \mathbf{w}\ $	$2/\ \mathbf{w}\ $	SV	Tr F1	Val F1	Val AUC
10^{-6}	24.75	209.41	0.010	8.9%	0.913	0.902	0.968
3×10^{-6}	8.25	75.66	0.026	9.3%	0.913	0.902	0.968
10^{-5}	2.47	25.03	0.080	10.4%	0.913	0.903	0.968
3×10^{-5}	0.82	10.85	0.184	12.9%	0.913	0.903	0.969
10^{-4}	0.25	6.04	0.331	17.2%	0.914	0.904	0.969
3×10^{-4}	0.08	4.62	0.433	20.3%	0.914	0.904	0.969
10^{-3}	0.02	3.81	0.525	22.6%	0.913	0.903	0.969
3×10^{-3}	0.01	2.98	0.670	24.7%	0.912	0.904	0.968
10^{-2}	0.00	2.15	0.929	28.0%	0.911	0.903	0.967
3×10^{-2}	0.00	1.62	1.233	33.1%	0.909	0.900	0.965
10^{-1}	0.00	1.13	1.766	41.8%	0.898	0.889	0.961

($C = 1/(n\lambda)$, LIBLINEAR [8]), both behind a StandardScaler inside a Pipeline. Our RFF SVM vs. the exact-kernel SVC with the same γ and C . A constant predictor (training median) anchors Task A. All randomness is seeded (SEED=42).

VI. RESULTS

A. Tree: criteria, depth and the bias–variance trade-off

Fig. 2 shows the classic picture. Training error keeps falling with depth, while validation error bottoms out and then rises again as the tree starts fitting noise. An unpruned regression tree reaches a training RMSE of 0.025 at depth 25, but its validation RMSE is 0.254; the best unpruned depth is 12. Gini and entropy behave almost identically on Task B: their best validation F1 differ by 0.001 (Gini 0.924, entropy 0.922). Larger leaves regularise a fully grown tree, and min_impurity_decrease prunes it: at 10^{-2} the tree shrinks from 3,125 to 11 nodes (Fig. 9). On validation we select **gini, depth 15, min_samples_leaf 1** for Task B and **depth 21, min_samples_leaf 20, min_impurity_decrease 0** for Task A (validation RMSE 0.229). In the learning curve (appendix) the train–validation gap goes from 0.050 to 0.025 as data grows, and validation error is still falling at the full training size, so more data would still help.

B. SVM: regularisation, margin and convergence

Table 3 and Fig. 3 (top) trace the effect of λ . A larger λ shrinks $\|\mathbf{w}\|$, which widens the margin $2/\|\mathbf{w}\|$ and makes more training points fall inside it as support vectors, trading training fit for a simpler boundary. Validation selects $\lambda = 3 \times 10^{-4}$ ($C = 0.082$). At this setting the final model has $\|\mathbf{w}\| = 4.479$, a margin of 0.447, and 10,197 support vectors (20.8% of the training rows).

Fig. 3 (middle) confirms convergence: J falls from 1.000 at $\mathbf{w} = 0$ to 0.2221 and changes by 0.18% over the last five epochs. Small λ converges slowly, as the $1/(\lambda T)$ rate predicts, which is why averaging matters. With the same epoch budget, mini-batches of 32 make as much progress per epoch as $k = 1$ at a fraction of the wall time (batch 1, last iterate: 26.91s; batch 32, last iterate: 1.81s; batch 32, averaged: 1.96s; full batch: 1.48s), while full-batch descent, which takes only one step per epoch, lags far behind. On the full training set, our averaged Pegasos ends at $J = 0.2238$, against LIBLINEAR’s exact optimum $J = 0.2224$ (gap 0.61%), and the two agree on 99.6% of test predictions.

For the RFF extension, validation selects $\gamma = 0.003$ and $\lambda = 10^{-6}$ (Fig. 3, bottom). Validation picks the smallest bandwidth γ and the weakest regularisation λ in the grid: the RBF kernel prefers to be nearly linear on these features, so the extension adds little over the linear SVM. The kernel-approximation error falls with a fitted log–log slope of -0.51 (theory: -0.5 , i.e. $1/\sqrt{D}$), and validation AUC saturates well before the largest D .

C. Benchmark against scikit-learn

Tables 4–6 are the main results. *Task A*: our tree’s test RMSE is 0.229 (CV 0.239 ± 0.004), scikit-learn’s is 0.229 (CV 0.239 ± 0.004), and the constant baseline’s is 0.628. The back-transformed error is 43,258 AZN MAE (MAPE 15.6%). The bonus ridge regression (0.237) shows how much of the signal is close to linear in log space: the single tree beats it. *Task B*: the best of our models is our decision tree (F1 0.926), while the best model overall is sklearn DecisionTreeClassifier (F1 0.928). The better of our two SVMs trails our tree, and the RFF kernel matches the linear SVM (test F1 +0.002). Averaged Pegasos is within 0.000 F1 of LIBLINEAR.

Structure. Table 7 compares the fitted trees node by node (src/tree_compare.py). At each node we check whether both libraries split the same training rows. If they do not, we classify the difference: an exact tie (two splits with equal gain; scikit-learn breaks ties by visiting features in random order), a *float32 divergence* (scikit-learn casts X to float32, which here merges 33,547 distinct float64 values; we re-run our own split search on the float32-rounded rows of that node and it returns exactly scikit-learn’s split), or an unexplained difference. We find 0 (Task A) and 0 (Task B) unexplained differences, 124 ties and 19 float32 divergences in Task A, and test-prediction agreement of 94.5% and 98.9%. The root splits are area_m2 ≤ 100.100 (ours) and log_area ≤ 4.606 (scikit-learn).

Speed. Our trees fit in 5.68 s vs. 0.671 s for scikit-learn’s Cython code. Our split search is asymptotically the same $O(dm \log m)$ per node, but it pays Python overhead per node and re-sorts at every node instead of reusing a presorted order. Pegasos costs $O(\text{epochs} \cdot nd)$ and fits in 2.50 s. The exact-kernel SVC is $O(n^2) - O(n^3)$ (10.53 s), the main practical argument for RFF.

D. Error analysis

What the models use. Our tree’s impurity-based importances (Fig. 4) are led by area_m2 (0.61), dist_center_km (0.21), lng (0.05), land_area_sot (0.04) for Task A and by area_m2 (0.53), dist_center_km (0.15), lng (0.08), lat (0.04) for Task B. Size and location dominate, which matches the EDA. In the linear SVM (Fig. 6), whose weights are comparable because the features are standardised, the largest positive weights are log_area, location=ag seher q., lat, lng and the most negative are dist_center_km, area_m2, location=bine q., district=suraxani r..

Where the regression fails. (i) *Luxury and bargain outliers*. The tree under-predicts the top price decile (mean log residual -0.127) and over-predicts the bottom one (0.136, Fig. 5, bottom left). Leaves average many listings, which pulls the extremes towards the mean; this is the variance-






Figure 2. Depth sweep with `min_samples_leaf=1` (dotted: train, solid: validation). Validation error is U-shaped, the bias–variance trade-off. The grey line marks the depth chosen on validation.



Top: effect of λ on (a) validation F1/AUC, (b) the margin width $2/\|w\|$, (c) the share of support vectors.



Middle: training objective per epoch for (a) different λ , (b) batch size and iterate averaging at the selected λ , (c) step-size schedules.




Figure 3. Our Pegasos SVM. **Bottom:** RFF extension: (a) bandwidth $\gamma \times \lambda$, (b) validation AUC vs. number of random features D , (c) mean absolute error of the kernel approximation.

Table 4. Task A, regression of $\log(\text{price})$. CV: 5-fold mean $\pm$ std on train+validation. Test: single held-out split with a 95% bootstrap CI. MAE (AZN) and MAPE after back-transformation. Best test RMSE in bold.

Model	CV RMSE (log)	CV R^2 (log)	Test RMSE (log) [95% CI]	Test R^2	Test MAE (AZN)	MAPE	Fit (s)
Constant baseline (train median)	—	—	0.628	-0.001	123,956	54.2%	0.00
Ours: decision tree	0.239 $\pm$ 0.004	0.860 $\pm$ 0.004	0.229 [0.222, 0.237]	0.866	43,258	15.6%	5.68
sklearn DecisionTreeRegressor	0.239 $\pm$ 0.004	0.860 $\pm$ 0.004	0.229 [0.222, 0.237]	0.866	43,349	15.6%	0.67
Ours: ridge (bonus)	0.249 $\pm$ 0.000	0.849 $\pm$ 0.002	0.237 [0.230, 0.244]	0.857	49,115	17.0%	0.03
sklearn Ridge	0.249 $\pm$ 0.000	0.849 $\pm$ 0.002	0.237 [0.230, 0.244]	0.857	49,115	17.0%	0.04

Table 5. Task B, price tier (premium = positive class). Same protocol as Table 4. Best test F1 in bold. † trained on a 20,000-row subsample ($O(n^2)$ kernel solver).

Model	CV F1	CV ROC-AUC	Acc	Prec	Rec	Test F1 [95% CI]	AUC	PR-AUC	Fit (s)
Ours: decision tree	0.923 $\pm$ 0.003	0.909 $\pm$ 0.004	0.927	0.932	0.921	0.926 [0.921, 0.932]	0.921	0.904	15.17
sklearn DecisionTreeClassifier	0.922 $\pm$ 0.003	0.909 $\pm$ 0.004	0.929	0.933	0.924	0.928 [0.922, 0.934]	0.924	0.908	0.86
Ours: linear SVM (Pegasos)	0.911 $\pm$ 0.004	0.971 $\pm$ 0.001	0.911	0.907	0.916	0.912 [0.906, 0.918]	0.972	0.972	2.50
sklearn SGDClassifier (hinge)	0.908 $\pm$ 0.004	0.970 $\pm$ 0.001	0.908	0.902	0.915	0.909 [0.902, 0.915]	0.970	0.972	1.27
sklearn LinearSVC	0.911 $\pm$ 0.004	0.971 $\pm$ 0.001	0.911	0.906	0.917	0.911 [0.905, 0.918]	0.972	0.972	5.05
Ours: RBF SVM (RFF + Pegasos)	0.914 $\pm$ 0.003	0.974 $\pm$ 0.001	0.913	0.911	0.916	0.914 [0.907, 0.920]	0.974	0.974	13.86
sklearn SVC (RBF) †	0.917 $\pm$ 0.002	0.976 $\pm$ 0.001	0.918	0.915	0.921	0.918 [0.912, 0.924]	0.976	0.976	10.53

Table 6. Headline results on both tasks (Task A: RMSE of log price, lower is better; Task B: F1, higher is better).

Model	A: RMSE (CV)	A: test	B: F1 (CV)	B: test
Ours: decision tree	0.239 $\pm$ 0.004	0.229	0.923 $\pm$ 0.003	0.926
Ours: linear SVM (Pegasos)	n/a	n/a	0.911 $\pm$ 0.004	0.912
Ours: RBF SVM (RFF + Pegasos)	n/a	n/a	0.914 $\pm$ 0.003	0.914
Ours: ridge (bonus)	0.249 $\pm$ 0.000	0.237	n/a	n/a
sklearn DecisionTree*	0.239 $\pm$ 0.004	0.229	0.922 $\pm$ 0.003	0.928
sklearn LinearSVC	n/a	n/a	0.911 $\pm$ 0.004	0.911
sklearn SVC (RBF)	n/a	n/a	0.917 $\pm$ 0.002	0.918

Table 7. Our tree vs. scikit-learn, node by node (final models; depth and leaves as ours / sklearn). Same: identical split. Ties: equal-gain alternatives. f32: explained by scikit-learn’s float32 cast. Other: unexplained (0 for a correct implementation).

Task	Depth	Leaves	Same	Ties	f32	Other	Agree
A (MSE)	21 / 21	1,869 / 1,872	1,375	124	19	0	94.5%
B (gini)	15 / 15	1,898 / 1,906	1,079	400	9	0	98.9%

reduction side of the bias–variance trade-off. Table 8 lists the largest errors. (ii) *Sparse districts* (Fig. 5, top middle). The test MAE is 0.382 for listings whose district has <20 training examples, against 0.127 for 500+: rare districts fall into the “other” column and fall back on coordinates. (iii) *Categories*. The error is highest for heyet evi/bag evi (0.299) and lowest for kohne tikili (0.114). 4.9% of test listings are off by more than 50% in price.

Where the classifier fails. Tier errors concentrate at the boundary. 42% of the tree’s misclassifications (SVM: 49%) are listings within $\pm 10\%$ of the threshold price, a band that holds only 11% of the test rows. The error rate falls from 33.3% closest to the threshold to 1.8% far from it (Fig. 5, bottom right). Many of these “errors” are therefore near-ties that no feature could resolve: a listing just above the median is not meaningfully more premium than one just below it.

E. Bonus: ridge, ensembles of our trees, clustering

Ridge (normal equations, $\alpha = 100$ chosen on validation) matches scikit-learn’s Ridge to machine precision. *Ensem-*

Table 8. Largest Task A test errors of our tree.

True (AZN)	Pred. (AZN)	m ²	Rooms	Category	Location
440,000	58,006	100	4	heyet evi/bag evi	bilgeh q.
818,000	128,424	450	8	yeni tikili	bineqedi r.
649,000	128,424	320	7	heyet evi/bag evi	bineqedi r.
1,300,000	265,663	300	11	heyet evi/bag evi	novxani q.
38,000	179,423	90	2	yeni tikili	yeni yasamal q.
20,000	91,588	90	3	heyet evi/bag evi	binc q.

Table 9. Bonus models on Task A (test split).

Task A model (test)	RMSE (log)	R^2	Fit (s)
Ours: single decision tree	0.229	0.866	—
Ours: ridge (normal equations)	0.237	0.857	—
Ours: random forest	0.180	0.918	250.7
sklearn RandomForest	0.181	0.917	8.7
Ours: gradient boosting	0.204	0.894	449.5
sklearn GradientBoosting	0.206	0.893	39.2

bles (Fig. 11, Table 9): a random forest of 40 of our trees (bootstrap + a third of the features per split) lowers the test RMSE: 0.180 against 0.229 for a single tree (scikit-learn’s forest: 0.181). Least-squares gradient boosting with 120 depth-4 trees reaches 0.204 (scikit-learn: 0.206). The boosting curve falls steadily as learners are added. For Task B our forest reaches F1 0.945 (scikit-learn 0.944). *Clustering*: k-means++ ($k = 12$) on the projected coordinates recovers neighbourhoods (Fig. 12). Per-cluster test MAE ranges from 0.119 to 0.546 (cluster 4). Adding cluster indicators to the tree does not lower the validation RMSE (0.229 $\rightarrow$ 0.229). PCA needs 71 components for 90% of the variance; the first two hold only 10.6%.

VII. DISCUSSION AND LIMITATIONS

Bias and variance. A single deep tree has low bias and high variance. Validation-chosen depth and leaf size cut the variance, and averaging many decorrelated trees (the forest) cuts it further, which is why our forest beats the single tree. The linear SVM has the opposite profile: it cannot express interactions such as “large and central”, and RFF buys that




Figure 4. Top impurity-decrease (MDI) importances of our final trees.




Figure 5. Error analysis on the test split. Top: Task A MAE of log(price) by category, by number of training listings in the district, and by true-price decile (row counts in brackets). Bottom left: mean residual per price decile; the tree pulls the extremes towards the mean. Bottom right: Task B error rate vs. distance from the threshold. A map of the residuals is in the appendix (Fig. 8).




Figure 6. Largest weights of the final linear SVM (standardised features; positive values push towards premium).

flexibility at the price of tuning γ .

Honest limitations. (1) Asking prices are not transaction prices, and re-posted listings with edited prices can survive de-duplication. (2) Cleaning thresholds and the residential scope are judgement calls. They are fixed and documented but not tuned, and filtering on price per m² removes some genuine extremes along with typos. (3) One six-week snapshot and one random split: CV quantifies sampling noise but not temporal drift. A time-based split would be the stricter test. (4) Task B is a median split of Task A’s target, so it inherits all of Task A’s noise and has an irreducible error band at the threshold. (5) Our tree is several times slower than scikit-learn’s, and Pegasos with very small λ needs many epochs. The RFF model is a Monte-Carlo approximation whose result depends on the random seed. (6) The description keywords are crude. A proper text model could add signal, but it would also need care to avoid prices written in the text.

VIII. CONCLUSION

Both from-scratch models behave as specified: the tree matches scikit-learn’s split at every node up to ties and float32 effects, and Pegasos reaches LIBLINEAR’s objective up to 0.61%. Both are competitive on real, messy data. Most of the work, and most of the error, lies in the data: currency quirks, duplicates, sparse districts and the luxury tail. Next steps are a time-based split, district-level target encoding learned inside the folds, quantile losses for the tails, and histogram-based split finding to close the speed gap.

CONTRIBUTIONS

See `contribution_report.md` in the repository root for who did what.

TOOLS AND ACKNOWLEDGEMENTS

Data: the Kaggle *Bina.az sale project* dataset [13]; we do not redistribute it. scikit-learn [9] is used only for the base-lines, the bonus comparisons and the unit tests that check

our metrics and models. **AI assistance:** an AI coding assistant (Claude Code, Anthropic) was used to draft and refactor parts of the code, tests and this report. All of it was reviewed, run and checked by the team, and every number in this report is regenerated by `python -m src.run_all`.

REFERENCES

- [1] S. Shalev-Shwartz, Y. Singer, N. Srebro, and A. Cotter, “Pegasos: primal estimated sub-gradient solver for SVM,” *Math. Programming*, vol. 127, no. 1, pp. 3–30, 2011.
- [2] L. Breiman, J. Friedman, R. Olshen, and C. Stone, *Classification and Regression Trees*. Wadsworth, 1984.
- [3] J. R. Quinlan, “Induction of decision trees,” *Machine Learning*, vol. 1, no. 1, pp. 81–106, 1986.
- [4] C. Cortes and V. Vapnik, “Support-vector networks,” *Machine Learning*, vol. 20, no. 3, pp. 273–297, 1995.
- [5] A. Rahimi and B. Recht, “Random features for large-scale kernel machines,” in *Proc. NeurIPS*, 2007, pp. 1177–1184.
- [6] S. Lacoste-Julien, M. Schmidt, and F. Bach, “A simpler approach to obtaining an $O(1/t)$ convergence rate for the projected stochastic subgradient method,” arXiv:1212.2002, 2012.
- [7] L. Bottou, “Large-scale machine learning with stochastic gradient descent,” in *Proc. COMPSTAT*, 2010, pp. 177–186.
- [8] R.-E. Fan, K.-W. Chang, C.-J. Hsieh, X.-R. Wang, and C.-J. Lin, “LIBLINEAR: A library for large linear classification,” *JMLR*, vol. 9, pp. 1871–1874, 2008.
- [9] F. Pedregosa *et al.*, “Scikit-learn: Machine learning in Python,” *JMLR*, vol. 12, pp. 2825–2830, 2011.
- [10] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [11] J. H. Friedman, “Greedy function approximation: a gradient boosting machine,” *Ann. Statist.*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [12] D. Arthur and S. Vassilvitskii, “k-means++: the advantages of careful seeding,” in *Proc. ACM-SIAM SODA*, 2007, pp. 1027–1035.
- [13] sehriyarmemmedli, “Bina.az sale project,” Kaggle dataset. <https://www.kaggle.com/datasets/sehriyarmemmedli/binaaz-sale-project>

APPENDIX: ADDITIONAL FIGURES




Figure 7. Location effect (median price per m² by district, training listings in brackets), premium share by category, and log price by room count (training split).




Figure 8. Test residuals of our Task A tree across Baku (red: over-predicted, blue: under-predicted).




Figure 9. Top: validation RMSE vs. `min_samples_leaf` for a fully grown tree, and vs. `min_impurity_decrease` at the selected depth. Bottom: learning curves of the selected trees.




Figure 10. Test-set ROC and precision–recall curves, confusion matrices, and predicted vs. actual log price.




Figure 11. Bonus: test RMSE as learners are added to our random forest (a) and our gradient boosting (b), with scikit-learn's final value and our single tree as references.




Figure 12. Bonus: k-means elbow and geographic clusters coloured by median log price; PCA explained variance and the two tiers in the first two components.